

ThoughtWorks®

KUBERNETES

Rise of the Containers Workshop



“I predict this technology will be too good to resist, those who are not participating today will change their mind later.”

- Jim Zemlin, Linux Foundation

Brief History

- Greek for “helmsman” or pilot
- First announced in mid-2014, as an all-Google project
- In mid-2015, Google + Linux Foundation came together to form CNCF

What is it?

“Kubernetes is a portable, extensible open-source platform for managing containerized workloads and services”

Ansible

Configuration

Works with hosts directly

Provision a system with the required config

On-demand system



Kubernetes


Orchestration

Works with containers

Deploy ready-made images to infra

Live, realtime system







Kubernetes Architecture

Getting perspective at 30000 feet

High-level architecture



Master Components



kube-apiserver

- Frontend to the “control plane”
- Exposes a REST API to interact with kubernetes
- Often mistaken for being the master




etcd

- Distributed, key-value store
- Used to store all cluster data
- The only stateful component in kubernetes
- “Source of Truth”



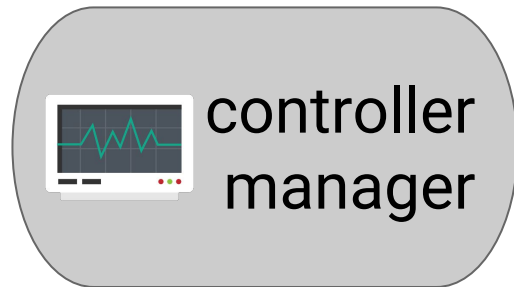
kube-scheduler

- Assigns a node for newly created pods
- Various options to choose from for affinity
- Continuously keeps watching store for new pods

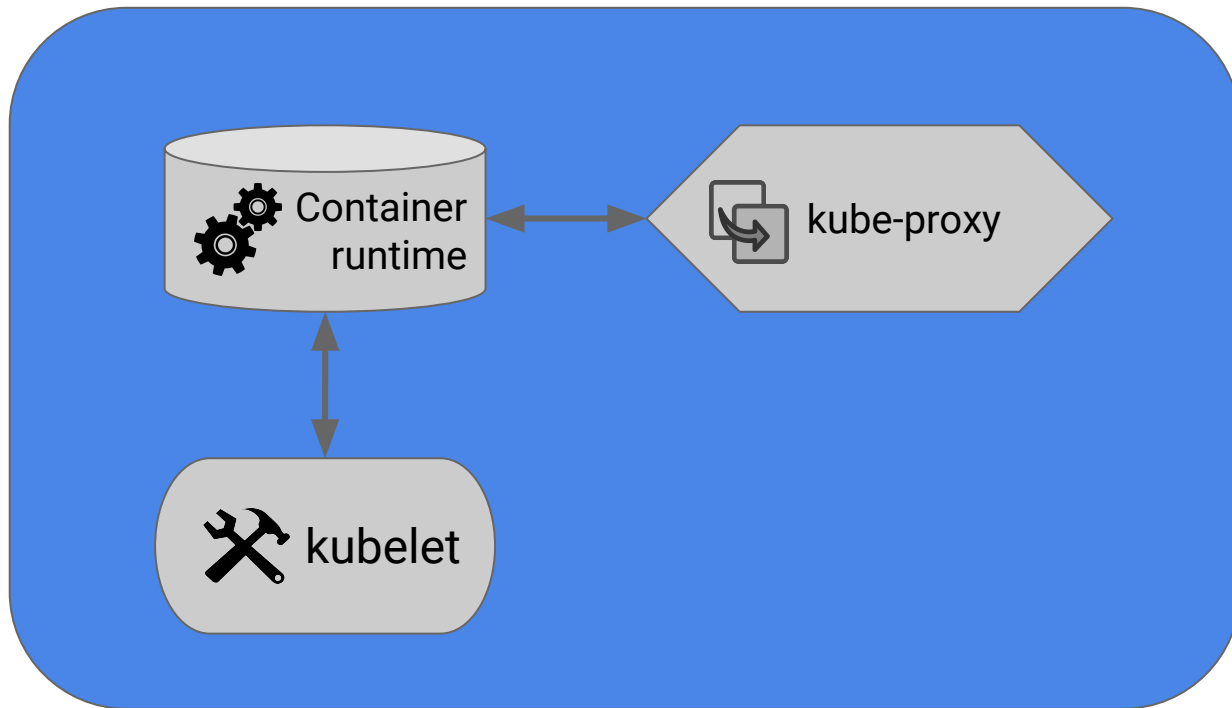


kube-controller-manager

- Runs “controllers” in kubernetes
- Controllers are daemons that ensure the desired state is achieved in the cluster
- Examples of controller:
 - Node Controller
 - Replication Controller
 - Endpoints Controller



Node Components



kubelet


- Main agent on the node, often mistaken for the node itself
- Keeps watching the apiserver for work
- Instantiates pods
- Reports to master



kubelet

Container Runtime

- Deals with the container abstraction
- Pull images, start/stop containers, etc
- Works with any OCI compliant container engine
- Usually docker, but also supports rkt

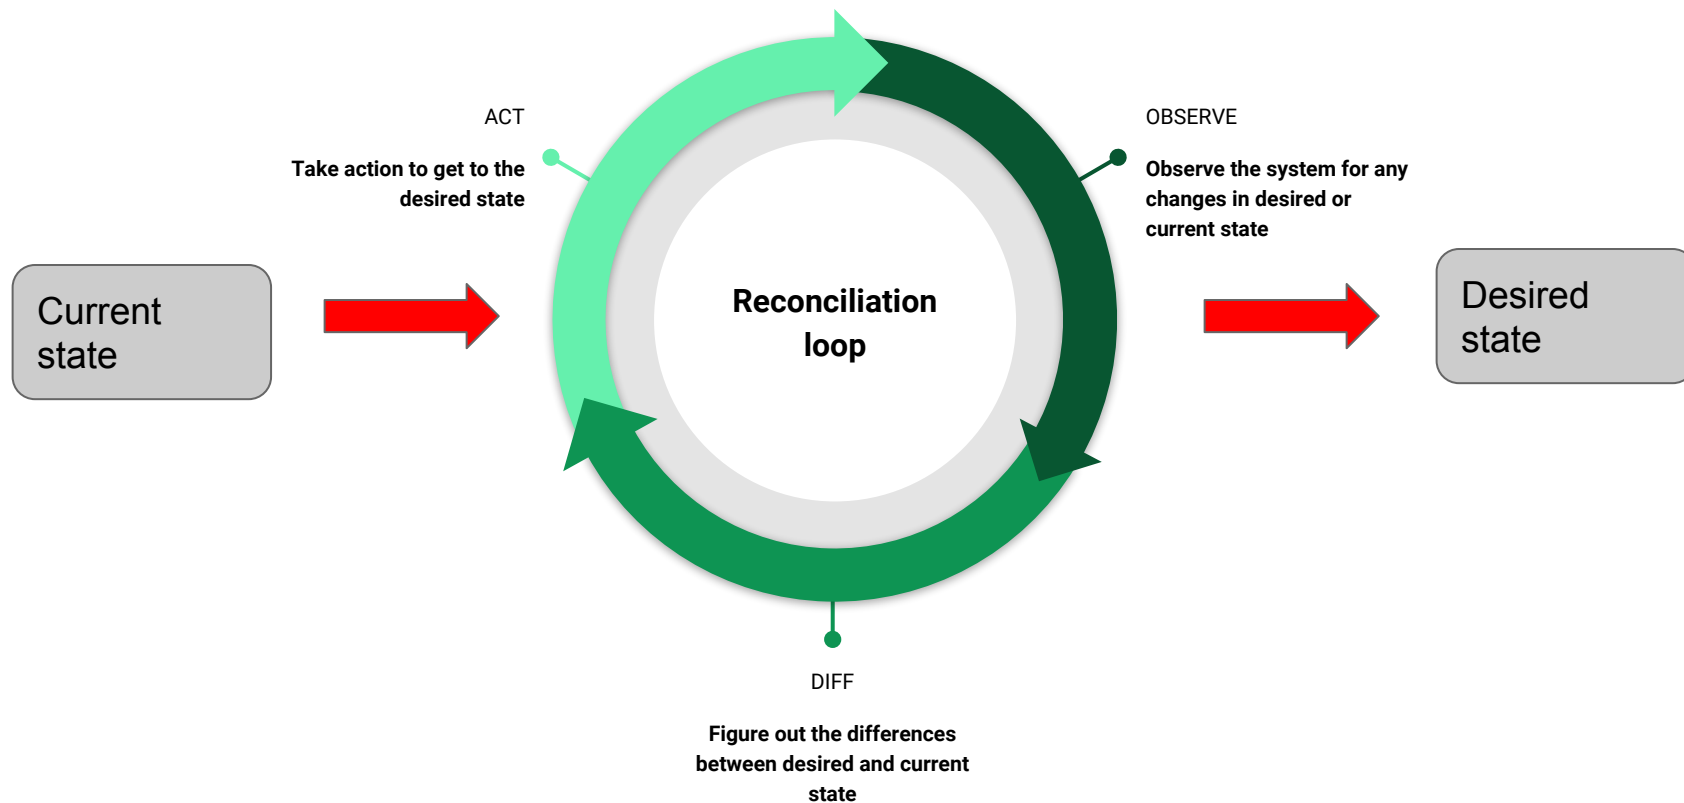


kube-proxy

- Deals with networking within the node
- Assigns an IP to each “pod”
- Is primarily used to maintain the “service” abstraction



Reconciliation





minikube

Hands on Minikube

*Let's kick things off with \$ **minikube start***

kubectl the command line

kubectl [command] [TYPE] [NAME] [flags]

Basic commands to inspect k8s


```
$ kubectl version
```

```
$ kubectl cluster-info
```

```
$ kubectl top
```

```
$ kubectl config view
```

kubectl config



kubectl

config context

```
$ kubectl config get-contexts
```

```
$ kubectl config current-context
```

```
$ kubectl config use-context
```

```
$ kubectl --context [...]
```

kubectl namespaces

```
$ kubectl get namespaces
```

```
$ kubectl get pods
```

```
$ kubectl get pods \  
    --namespace kube-system
```

kubectl

running applications

```
$ kubectl run hello-world \  
    --image=gcr.io/google-samples/node-hello:1.0 \  
    --port=8080
```

```
$ kubectl get pods
```

```
$ kubectl describe pod <podName>
```


kubectl

access applications

```
$ kubectl expose deployment  
hello-world \  
--type=NodePort \  
--name=hello-world-service
```

```
$ kubectl describe services \  
hello-world-service
```

```
$ curl http://<minikube-ip>:<nodePort>
```

Quick exercise!

Run the **metadata** service using **kubectl**

NOTE:

- Point local docker client to minikube's docker:
`$ eval $(minikube docker-env)`
- Build the metadata image again before attempting to run it