

ThoughtWorks®

STORAGE & VOLUMES

Rise of the Containers Workshop



Topics

1. Volumes
2. Persistent Volumes
3. Persistent Volume Claims
4. Storage Classes

Containers are ephemeral



Volumes


- Just a directory, accessible to all containers in pod
- Lasts as long as a pod
- The backing medium and contents depend on type of the volume
- Some important volume types:
 - `emptyDir`
 - `hostPath`
 - `awsElasticBlockStore`
 - `gcePersistentDisk`
 - `nfs`
 - `persistentVolumeClaim`

The ***emptyDir*** volume type


- Created at pod creation time and accessible by all containers in the pod
- Stored on the backing node

```
apiVersion: v1
kind: Pod
metadata:
  name: empty-dir-pod
spec:
  containers:
    - image: gcr.io/google-samples/node-hello:1.0
      name: empty-dir-container
      volumeMounts:
        - mountPath: /cache
          name: cache-volume
  volumes:
    - name: cache-volume
      emptyDir: {}
```


Why do we need better volumes?




We need volumes that are available across the cluster and reliable.



We need to de-couple provision of storage from consumption of storage.



The perfect way to represent a volume in k8s



Persistent Volumes (PV)

- Persistent Volumes are a kubernetes resource that represents a volume
- Administered and provisioned independently of other resources
- Can be accessed from kubectl:

```
$ kubectl get pv
```

```
kind: PersistentVolume
```

Persistent Volume Claims (PVC)

- PVCs are kubernetes resources that represents a claim to a persistent volume (big surprise)
- A claim is a request for access to a persistent volume
- Claims can be used as volumes for pods

```
kind: PersistentVolumeClaim
```


Analogy Question (2 Marks)

PersistentVolumeClaim:PersistentVolume::_____:

- 1) Pod:Service
- 2) Service:Pod
- 3) Pod:Node
- 4) Pod:Deployment

Storage Classes

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: slow
provisioner: kubernetes.io/aws-efs
parameters:
  type: io1
  zone: us-east-1d
  iopsPerGB: "10"
```



Hands On Time!

```
$ minikube ssh
```

```
$ mkdir -p /tmp/test
```

```
$ echo "Hello Storage!" > /tmp/test/index.html
```

Let's create a persistent volume

```
kind: PersistentVolume
apiVersion: v1
metadata:
  name: mypv
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 500Mi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Retain
  hostPath:
    path: "/tmp/test"
```

*Now, create the
corresponding
volume claim*

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: mypvc
spec:
  storageClassName: manual
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 500Mi
```

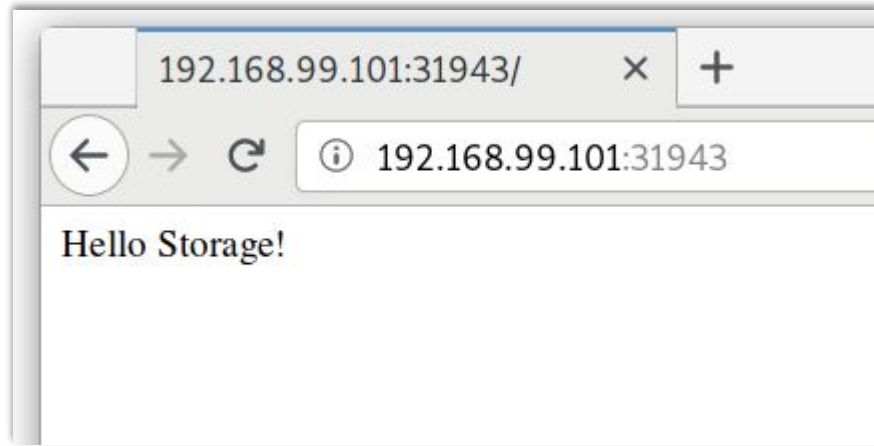
*Spin up a pod to
use the volume*

```
kind: Pod
apiVersion: v1
metadata:
  name: mywebserver
  labels:
    type: mywebserver
spec:
  volumes:
    - name: public
      persistentVolumeClaim:
        claimName: mypvc
  containers:
    - name: mywebserver
      image: nginx
      ports:
        - containerPort: 80
          name: "http-server"
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: public
```


*Create a service to
expose the app*

```
kind: Service
apiVersion: v1
metadata:
  name: mywebservice
spec:
  selector:
    type: mywebserver
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
      type: NodePort
```

*Access the
application in the
browser*



Exercise time!




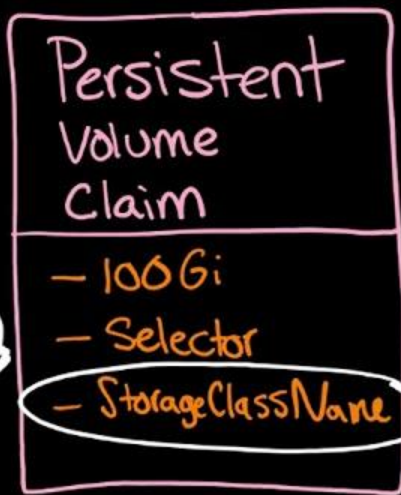
*Run the mongo database with the
persistent volume we created*

**P.S. The mount path for mongo should be
/data/db**

ThoughtWorks®



PVC



Persistent Volume Framework

